

# Java Threads - Synchronization

---

Università di Modena e Reggio Emilia

*Prof. Nicola Bicocchi (nicola.bicocchi@unimore.it)*



# Synchronization

---

- What happens when two different threads are accessing the same data ?
- Imagine that two people each have ATM cards, but both cards are linked to only one account

```
class Account {  
    private int balance;  
    public int getBalance() {return balance;}  
    public void withdraw(int amount) {  
        balance = balance - amount;  
    }  
}
```



# Synchronization

---

- Before one of them makes a withdrawal, first check the balance to be certain there's enough to cover the withdrawal
  - Check the balance.
  - If there's enough in the account, make the withdrawal
- What happens if something separates step 1 from step 2 ?



# Synchronization

---



# Synchronization

---

- Marge checks the balance and there is enough
- Before she withdraws money, Homer checks the balance and also sees that there's enough for his withdrawal. He is seeing the account balance before Marge actually debits the account...
- Both Marge and Homer believe there's enough to make their withdrawals !
- If Marge makes her withdrawal...
  - ...there isn't enough in the account for Homer's withdrawal
  - ... but he thinks there is since when he checked, there was enough!



# Synchronization

---



# Race condition

---

- A problem happening whenever:
  - Many threads can access the same resource (typically an object's instance variable)
  - This can produce corrupted data if one thread "*races in*" too quickly before an operation has completed.



# Preventing Race Conditions

---

- We must guarantee that the two steps of the withdrawal are NEVER split apart.
- It must be an atomic operation:
  - It is completed before any other thread code that acts on the same data
  - ...regardless of the number of actual instructions





# Preventing Race Conditions

---

- You can't guarantee that a single thread will stay running during the atomic operation.
- But even if the thread running the atomic operation moves in and out of the running state, no other running thread should be able to act on the same data.
- How to protect the data:
  - Mark the variables private
  - Synchronize the code that modifies the variables



# Synchronized

---

- The modifier synchronized can be applied to a method or a code block
- It locks\* a code block: **only one thread can access**

Remember: the simplest lock implementation is a semaphore (provided by OSs).



# Synchronization and Locks

---

- Every object in Java has one built-in lock
- Enter a synchronized non-static method => get the lock of the current object code we're executing.
- If one thread got the lock, other threads have to **wait** to enter the synchronized code until the lock has been released (thread exits the **synchronized** method)
- Not all methods in a class need to be synchronized.
- Once a thread gets the lock on an object, no other thread can enter **any** of the synchronized methods in that class (for that object).



# Synchronization and Locks

---

- Multiple threads can still access the class's non-synchronized methods
  - Methods that don't access the data to be protected, don't need to be synchronized
  - Threads going to sleep, don't release locks
- A thread can acquire more than one lock, e.g.
  - a thread can enter a synchronized method, then immediately invoke a synchronized method on another object



# Synchronization and Locks

---

```
public synchronized void doStuff() {  
    System.out.println("synchronized");  
}
```

Is equivalent to

```
public void doStuff() {  
    synchronized(this) {  
        System.out.println("synchronized");  
    }  
}
```



# Synchronization and Locks

---

```
public static synchronized int getCount() {  
    return count;  
}
```

Is equivalent to

```
public static int getCount() {  
    synchronized(MyClass.class) {  
        return count;  
    }  
}
```



# When Do I Need To Synchronize?

---

- Two threads executing the same method at the same time may:
  - use different copies of local vars => no problem
  - access fields that contain shared data
- To make a **thread-safe** class:
  - methods that access changeable fields need to be synchronized.
  - access to static fields should be done from static synchronized methods.
  - access to non-static fields should be done from non-static synchronized methods



# Synchronize the code yourself

---

```
class NameDropper implements Runnable{
    public void run() {
        String name = nl.removeFirst();
        System.out.println(name);
    }
}

public static void main(String[] args) {
    fNameList nl = new NameList();
    nl.add("Jacob");
    Thread t1 = new Thread(new NameDropper()); t1.start();
    Thread t2 = new Thread(new NameDropper()); t2.start();
}
```





# Synchronize the code yourself

---

```
public class NameList {  
    private List names = new LinkedList();  
  
    public synchronized void add(String name) {  
        names.add(name);  
    }  
  
    public synchronized String removeFirst() {  
        if (names.size() > 0)  
            return (String) names.remove(0);  
        else  
            return null;  
    }  
}
```



# Collections.synchronized\*

---

```
public class NameList {  
    private List names =  
        Collections.synchronizedList(new LinkedList());  
  
    public void add(String name) {  
        names.add(name);  
    }  
  
    public String removeFirst() {  
        if (names.size() > 0)  
            return (String) names.remove(0);  
        else return null;  
    }  
}
```



# Collections.synchronized\*

---

```
public static <T> List<T> synchronizedList(List<T> list)
```

Returns a synchronized (thread-safe) list backed by the specified list. In order to guarantee serial access, it is critical that all access to the backing list is accomplished through the returned list.

**It is imperative that the user manually synchronize on the returned list when iterating over it:**

```
List list = Collections.synchronizedList(new ArrayList());  
...  
synchronized (list) {  
    Iterator i = list.iterator(); // Must be in synchronized block  
    while (i.hasNext())  
        foo(i.next());  
}
```

Failure to follow this advice may result in non-deterministic behavior.



# Deadlock!

---

- Deadlock occurs when two threads are blocked, with each waiting for the other's lock.
  - Neither can run until the other gives up its lock, so they wait forever
- Poor design can lead to deadlock
  - It is hard to debug code to avoid deadlock
  - Model checking could be a solution (problem: state space explosion)



# Deadlock!

```
public class DeadlockProducer {  
    private static class Resource { public int value; }  
    private Resource resourceA = new Resource();  
    private Resource resourceB = new Resource();  
    public int read() {  
        synchronized(resourceA) { // May deadlock here !  
            synchronized(resourceB) {  
                return resourceB.value + resourceA.value;  
            }  
        }  
    }  
    public void write(int a, int b) {  
        synchronized(resourceB) { // May deadlock here !  
            synchronized(resourceA) {  
                resourceA.value = a; resourceB.value = b;  
            }  
        }  
    }  
}
```



# Livelock!

---

- A livelock happens when threads are actually running, but no work gets done
  - what is done by a thread is undone by another
- Example: each thread already holds one object and needs another that is held by the other thread.
  - What if each thread unlocks the object it owns and picks up the object unlocked by the other thread ?
  - They can run forever in lock-step!



# Starvation!

---

- Wait/notify primitives of the Java language do not guarantee **liveness** (=> starvation)
- When wait() method is called
  - thread releases the object lock prior to commencing to wait
  - and it must be reacquired before returning from the method, post notification



# Starvation!

---

- Once a thread releases the lock on an object (following the call to wait), it is placed in a object's wait-set
  - Implemented as a queue by most JVMs
  - When a notification happens, a new thread will be placed at the back of the queue
- By the time the notified thread actually gets the monitor, the condition for which it was notified may no longer be true ...
  - It will have to wait again
  - This can continue indefinitely => **Starvation**





# Synchronization in Object class

---

- void **wait()**
  - Causes current thread to wait until another thread invokes the **notify()** method or the **notifyAll()** method for this object.
- void **notify()**
  - Wakes up a single thread that is waiting on this object's lock.
- void **notifyAll()**
  - Wakes up all threads that are waiting on this object's lock.



# Mutual Exclusion

---

- A thread invokes wait() or notify() on a particular object, and the thread must currently hold the lock on that object
  - Called from within a synchronized context
- A thread owns a critical region when he has called wait() and it has not released the object yet (calling notify)
- A critical region is **nonsignaled** when is owned by a thread, **signaled** otherwise.



# Wait

---

- The wait() method lets a thread say:  
*“There's nothing for me to do now, so put me in your waiting pool and notify me when something happens that I care about.”*
- When calling wait() :
  - On a **signaled** object lock thread keeps executing
  - On a **nonsignaled** object lock thread is suspended



# Notify

---

- The `notify()` method send a signal to one of the threads that are waiting in the same object's waiting pool.
- The `notify()` method CANNOT specify which waiting thread to notify.
- The method `notifyAll()` is similar but only it sends the signal to all of the threads waiting on the object.



# Recap: Threads Are Hard

---

- Synchronization:
  - Must coordinate access to shared data with locks.
  - Forget a lock? Corrupted data.
- Deadlock:
  - Circular dependencies among locks.
  - Each process waits for some other process: system hangs.



# Recap: Threads Are Hard

---

- Achieving good performance is hard:
  - Simple locking yields low concurrency.
  - Fine-grained locking increases complexity
  - O.S. limits performance (context switches)
- Threads not well supported:
  - Hard to port threaded code (PCs? Macs?).
  - Standard libraries not thread-safe.
- Hard to debug :
  - data dependencies
  - timing dependencies
  - Few debugging tools

